

Exam 1

Author: Gabriel Hofer

Course: CSC-410 Parallel Computing

Instructor: Dr. Karlsson

Due: October 19, 2020

How to Make the Project

Listing 1: make

```
7356111@linux09 CSC-410-Exam-1 >>make
gcc openmp.c -fopenmp -lm -o openmp
nvcc cuda.cu -o cuda
pdflatex exam1.tex
...
```

Functions and Program Structure

main calls either usage or range or correctness depending on the command line arguments.

Usage prints a Message to standard output about how to start the program

makeMatrix makes a random matrix. the probability of that there is an edge for any two vertices is equal to 0.25. We use the rand() C language function to “generate” random integers. we also set the seed value before any rand() calls.

serial is our implementation of Floyd’s algorithm without any parallelization. We use it to check the correctness of our parallelized functions.

Correctness The purpose of this function is to test whether our parallelized code is correct. We make the assumption that the function called serial (previously mentioned) is a correct implementation of Floyd’s algorithm. So, we compare the output of our parallelized function to the output of serial.

Range runs Floyd’s algorithm in parallel for a range of values of n. we iterate from small power of 2 to a greater power of 2.

printA simply prints the 2D array, with a tab separating each column.

Floyd CUDA

Listing 2: Floyd CUDA

```
--global-- void aux(int * dA, const int n, const int k){
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    if(index >= n*n) return;
    __syncthreads();
    int i = index / n, j = index % n;
    dA[i*n+j] = dA[i*n+j] < (dA[i*n+k]+dA[k*n+j]) ?
        dA[i*n+j] : dA[i*n+k]+dA[k*n+j];
    __syncthreads();
}

void floyd(int * dA, const int n){
    for(int k=0;k<n;k++){
        aux<<<(n*n+THREADS_PER_BLOCK)/(THREADS_PER_BLOCK),
            THREADS_PER_BLOCK>>>(dA,n,k);
        cudaDeviceSynchronize();
    }
}
```

Floyd OpenMP

Listing 3: Floyd OpenMP

```
void floyd(int * A, const int n){
    for(int k=0;k<n;k++)
        #pragma omp parallel for
        for(int i=0;i<n;i++)
            #pragma omp parallel for
            for(int j=0;j<n;j++)
                A[i*n+j] = A[i*n+j] < (A[i*n+k]+A[k*n+j]) ?
                    A[i*n+j] : A[i*n+k]+A[k*n+j];
}
```

In the OpenMP version of Floyd's algorithm, we insert two pragmas for the second and third nested for loops.

Testing and Verification

Listing 4: CUDA Correctness Testing

```
void correctness(const int low, const int high){
    for(int n = pow(2,low); n <= pow(2,high); n*=2){
        int * A = makeMatrix(n);
        int * B = (int *)malloc(n*n*sizeof(int));
        int Asize = n*n*sizeof(int);
        memcpy(B, A, Asize);
        serial(B,n);

        int * dA=NULL;
        cudaMalloc((void **)&dA, Asize);
        cudaMemcpy(dA, A, Asize, cudaMemcpyHostToDevice);
        floyd(dA,n);
        cudaMemcpy(A, dA, Asize, cudaMemcpyDeviceToHost);

        bool foundDiff=false;
        for(int i=0;i<n;i++)
            for(int j=0;j<n;j++)
                if(B[i*n+j]!=A[i*n+j]){
                    foundDiff=true;
                    return;
                }
        cudaFree(dA);
        free(A);
        free(B);
        cudaDeviceSynchronize();
        if(foundDiff){
            printf("FOUND_DIFFERENCE:(\n\n");
            return;
        }
        printf("SAME\n");
    }
    printf("ALL_SAME:(\n\n");
}
```

The correctness testing functions for CUDA and OpenMP are similar.

Usage Examples

Listing 5: CUDA Range Option

```
7356111@linux09 CSC-410-Exam-1 >>./cuda -r 1 12
2, 0.000023
4, 0.000031
8, 0.000059
16, 0.000113
32, 0.000224
64, 0.000483
128, 0.001060
256, 0.003196
512, 0.014511
1024, 0.113596
2048, 0.846598
4096, 6.665819
```

Listing 6: OpenMP Range Option

```
7356111@linux09 CSC-410-Exam-1 >>./openmp -r 1 12
2, 0.000549
4, 0.000271
8, 0.000282
16, 0.000430
32, 0.001399
64, 0.001146
128, 0.008198
256, 0.033451
512, 0.168379
1024, 1.088660
2048, 7.835085
4096, 60.018689
```

The two listings above outputted 12 records with 2 fields in each record. The first field in each record is n and the second field in each record is the time it took to run Floyd’s algorithm with an n by n matrix.

Charts and Analysis

Deliverables

- 1. Makefile
- 2. cuda.cu
- 3. openmp.c
- 4. exam1.pdf