

# Exam 1

Author: Gabriel Hofer

Course: CSC-410 Parallel Computing

Instructor: Dr. Karlsson

Due: October 19, 2020

## How to Make the Project

```
7356111@linux09 CSC-410-Exam-1 >>make
gcc openmp.c -fopenmp -lm -o openmp
nvcc cuda.cu -o cuda
```

## Usage Statement

```
Usage: ./cuda [-r low_power high_power] [-c low_power high_power]
-r: runs Floyd's algorithm in parallel on range of powers of two
-c: runs correctness tests on range of powers of two
```

## Usage Examples (Performance Tests for a Range of n)

```
7356111@linux09 CSC-410-Exam-1 >>./cuda -r 1 12
2, 0.000023
4, 0.000031
8, 0.000059
16, 0.000113
32, 0.000224
64, 0.000483
128, 0.001060
256, 0.003196
512, 0.014511
1024, 0.113596
2048, 0.846598
4096, 6.665819
```

```
7356111@linux09 CSC-410-Exam-1 >>./openmp -r 1 12
2, 0.000549
4, 0.000271
8, 0.000282
16, 0.000430
32, 0.001399
64, 0.001146
128, 0.008198
256, 0.033451
512, 0.168379
1024, 1.088660
2048, 7.835085
4096, 60.018689
```

## Data in a Table

n	CUDA (sec)	OpenMP (sec)
2	0.000023	0.000549
4	0.000031	0.000271
8	0.000059	0.000282
16	0.000113	0.000430
32	0.000224	0.001399
64	0.000483	0.001146
128	0.001060	0.008198
256	0.003196	0.033451
512	0.014511	0.168379
1024	0.113596	1.088660
2048	0.846598	7.835085
4096	6.665819	60.018689

## Usage Examples (Correctness Tests)

```
7356111@linux09 CSC-410-Exam-1 >>./cuda -c 1 10
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
ALL SAME:)
```

```
7356111@linux09 CSC-410-Exam-1 >>./openmp -c 1 10
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
SAME
ALL SAME:)
```

For both options -r and -c, two integers are required to follow. These two integers are the range for  $p$ . The two listings above show how to run cuda.cu and openmp.c from the command line using the -c option (which runs correctness tests). Each output line that prints "SAME" is a successful test. And, "ALL SAME:)" means that all tests passed. Each test is run on a different value of  $n$  calculated as follows:

$$n = 2^p$$

Where  $p$  is a number between 1 and 10 which were the low\_power and high\_power arguments passed to the program.

## Contents of exam1.tgz

1. Makefile
2. cuda.cu
3. openmp.c
4. exam1.pdf

## Functions

- **main** calls either usage or range or correctness depending on the command line arguments.
- **Usage** prints a Message to standard output about how to run the program.
- **makeMatrix** makes a random matrix. the probability of that there is an edge for any two vertices is equal to 0.25. We use the rand() C language function to “generate” random integers. we also set the seed value before any rand() calls.
- **serial** is our implementation of Floyd’s algorithm without any parallelization. We use it to check the correctness of our parallelized functions.
- **Correctness** tests whether our parallelized code is correct. We make the assumption that the function called serial is a correct implementation of Floyd’s algorithm. So, we compare the output of our parallelized function to the output of serial.
- **Range** runs Floyd’s algorithm in parallel for a range of values of n. we iterate from small power of 2 to a greater power of 2.
- **printA** simply prints the 2D array, with a tab separating each column.

## Files

- **cuda.cu** contains entire CUDA implementation of floyd’s algorithm.
- **openmp.c** constains entire OpenMP implementation of floyd’s algorithm.

## CUDA Floyd

Listing 1: CUDA Floyd

```
--global-- void aux(int * dA, const int n, const int k){
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    if(index >= n*n) return;
    __syncthreads();
    int i = index / n, j = index % n;
    dA[i*n+j] = dA[i*n+j] < (dA[i*n+k]+dA[k*n+j]) ?
        dA[i*n+j] : dA[i*n+k]+dA[k*n+j];
    __syncthreads();
}

void floyd(int * dA, const int n){
    for(int k=0;k<n;k++){
        aux<<<(n*n+THREADS_PER_BLOCK)/(THREADS_PER_BLOCK),
            THREADS_PER_BLOCK>>>(dA,n,k);
        cudaDeviceSynchronize();
    }
}
```

In CUDA, we use indexing to get rid of the two inner for loops in Floyd's algorithm.

## OpenMP Floyd

Listing 2: OpenMP Floyd

```
void floyd(int * A, const int n){
    for(int k=0;k<n;k++)
        #pragma omp parallel for
        for(int i=0;i<n;i++)
            #pragma omp parallel for
            for(int j=0;j<n;j++)
                A[i*n+j] = A[i*n+j] < (A[i*n+k]+A[k*n+j]) ?
                    A[i*n+j] : A[i*n+k]+A[k*n+j];
}
```

In the OpenMP version of Floyd's algorithm, we insert two pragmas for the second and third nested for loops.

## Serial (Non-Parallelized)

Listing 3: Serial Floyd

```
void serial(int * A, const int n){
    for(int k=0;k<n;k++)
        for(int i=0;i<n;i++)
            for(int j=0;j<n;j++)
                A[i*n+j] = A[i*n+j] < (A[i*n+k]+A[k*n+j]) ?
                    A[i*n+j] : A[i*n+k]+A[k*n+j];
}
```

## Testing and Verification

Listing 4: CUDA Correctness Testing

```
void correctness(const int low, const int high){
    for(int n = pow(2,low); n <= pow(2,high); n*=2){
        int * A = makeMatrix(n);
        int * B = (int *)malloc(n*n*sizeof(int));
        int Asize = n*n*sizeof(int);
        memcpy(B, A, Asize);
        serial(B,n);

        int * dA=NULL;
        cudaMalloc((void **)&dA, Asize);
        cudaMemcpy(dA, A, Asize, cudaMemcpyHostToDevice);
        floyd(dA,n);
        cudaMemcpy(A, dA, Asize, cudaMemcpyDeviceToHost);

        bool foundDiff=false;
        for(int i=0;i<n;i++)
            for(int j=0;j<n;j++)
                if(B[i*n+j]!=A[i*n+j]){
                    foundDiff=true;
                    return;
                }
        cudaFree(dA);
        free(A);
        free(B);
        cudaDeviceSynchronize();
        if(foundDiff){
            printf("FOUND_DIFFERENCE:(\n\n");
            return;
        }
        printf("SAME\n");
    }
    printf("ALL_SAME:(\n\n");
}
```

Experimental Serial Fraction

$$e = \frac{\frac{1}{\psi} - \frac{1}{p}}{1 - \frac{1}{p}}$$